Patricia Gschoßmann

Model Pruning

Standard Pruning

How to?

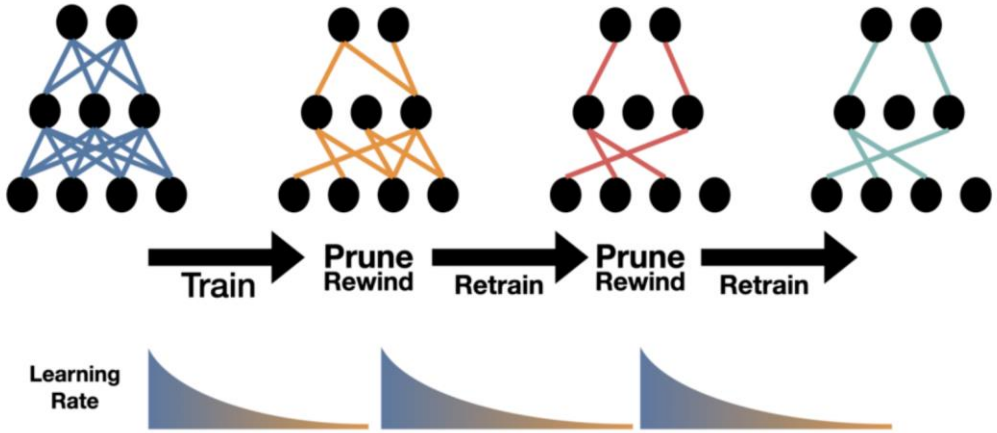

https://news.mit.edu/2020/foolproof-way-shrink-deep-learning-models-0430

Standard Pruning

How to?

How can you speed up the process?

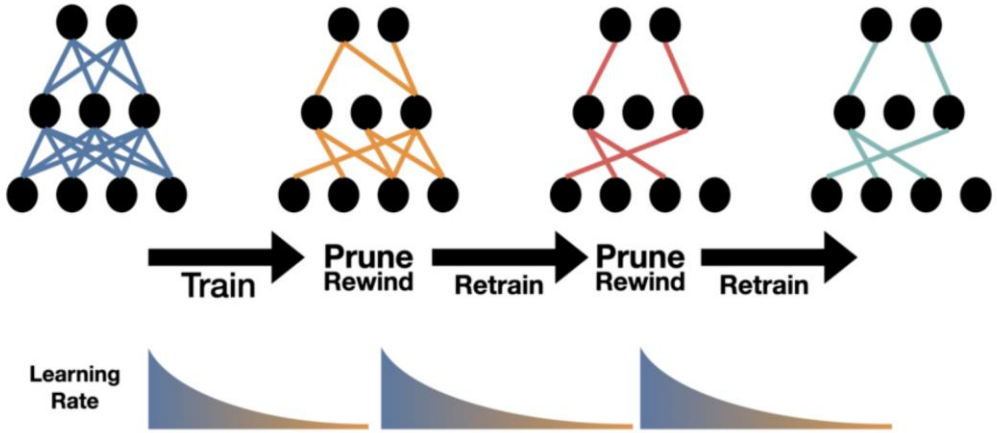

https://news.mit.edu/2020/foolproof-way-shrink-deep-learning-models-0430

New Approach

How to?

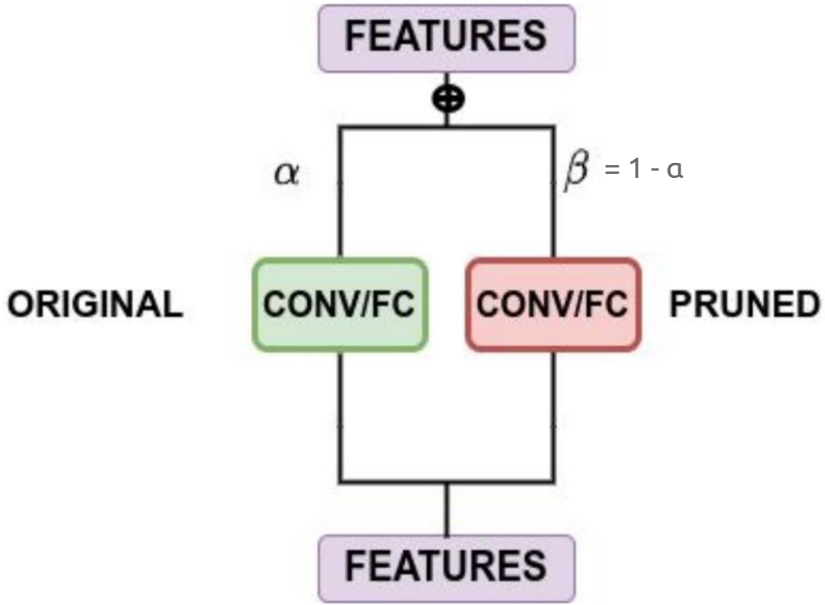

Learning for Self Driving Cars - Tutorial Week 7

New Approach

How to?

Advantage: Start at e.g. 60%

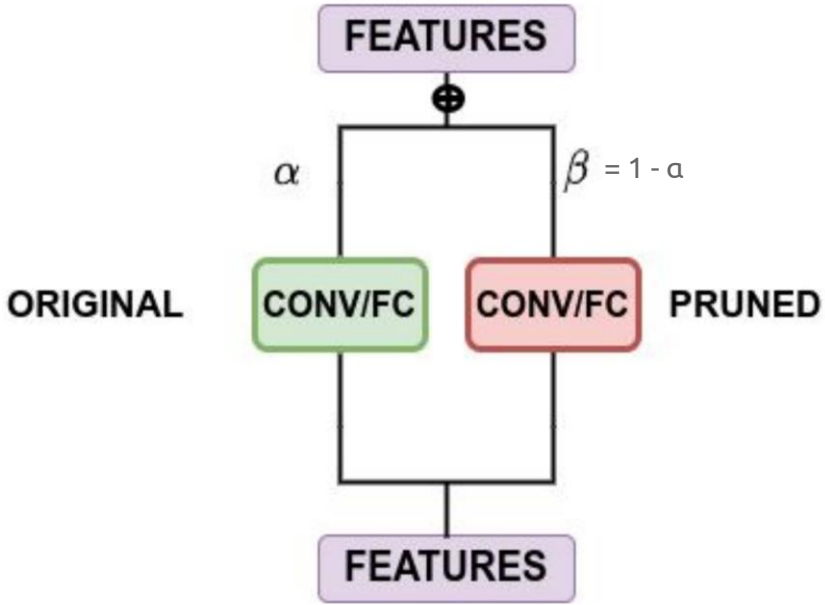

Learning for Self Driving Cars - Tutorial Week 7

Original model

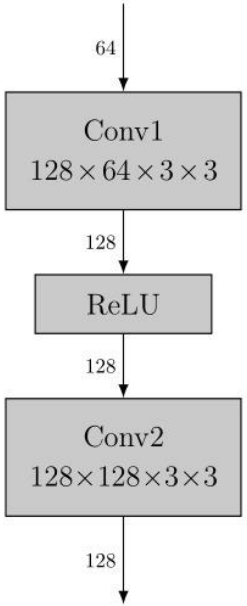

Original model

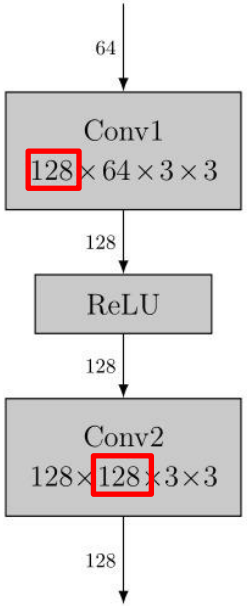

Original model

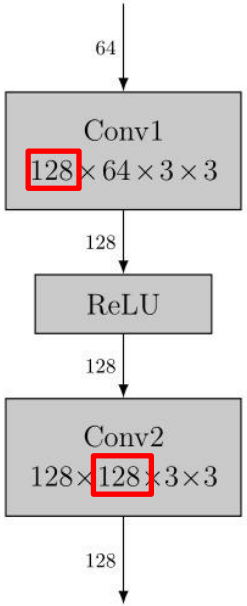

With parallel branches

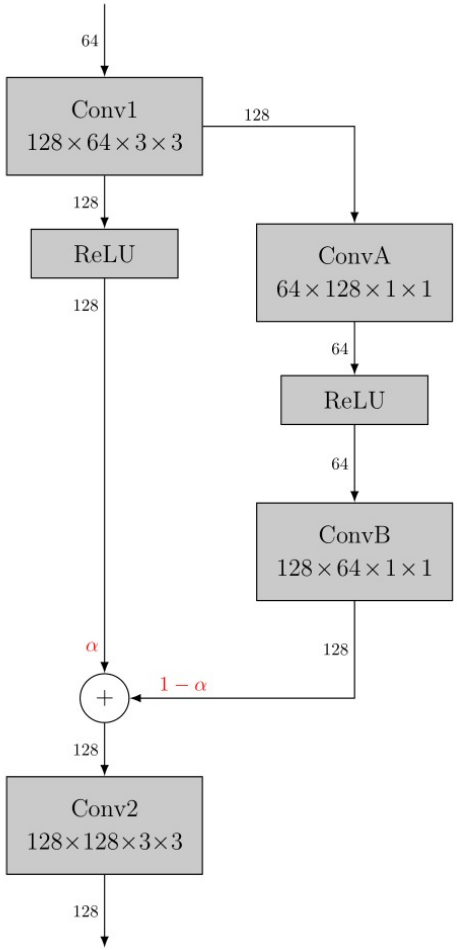

Original model

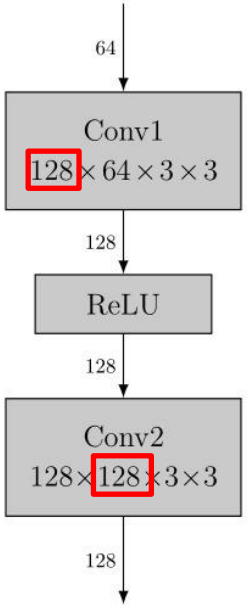

With parallel branches

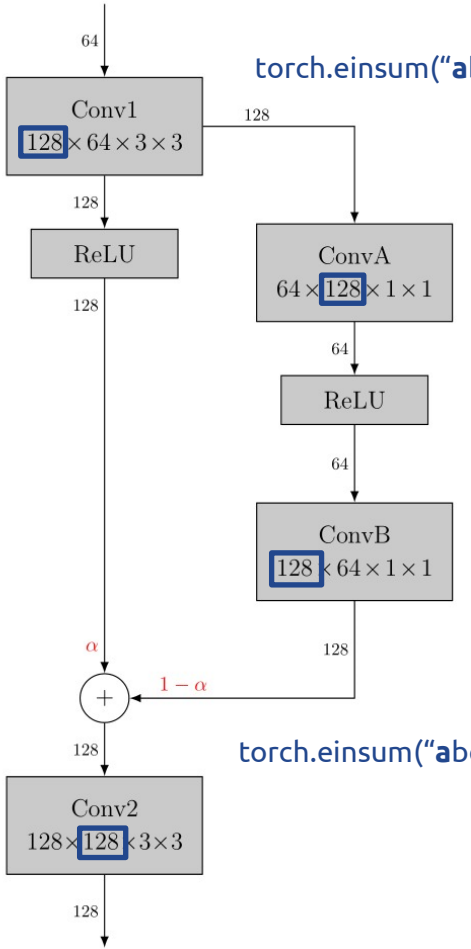

torch.einsum("**a**bcd, e**a**fg -> ebcd", conv1_weights, convA_weights)

torch.einsum("**a**bcd, e**a**fg -> ebfg", convB_weights, conv2_weights)

Original model

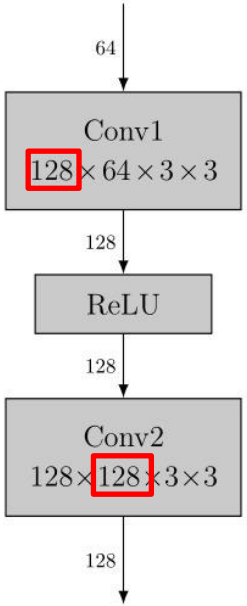

With parallel branches

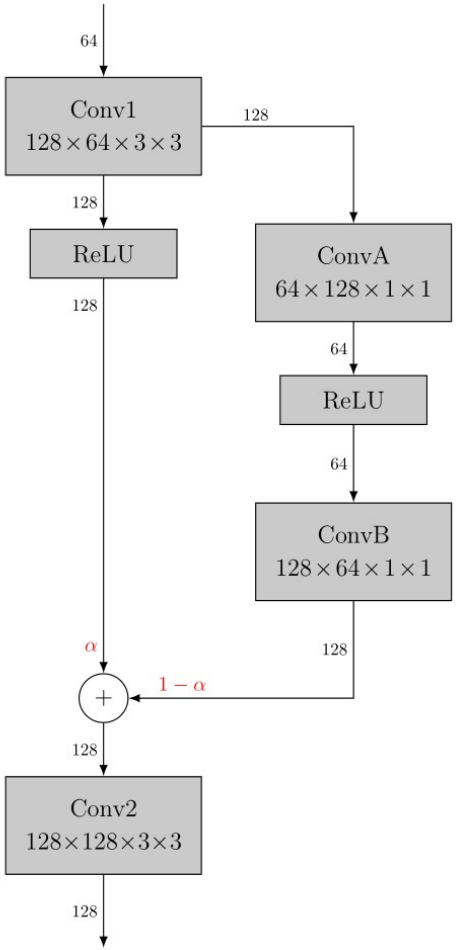

Pruned model

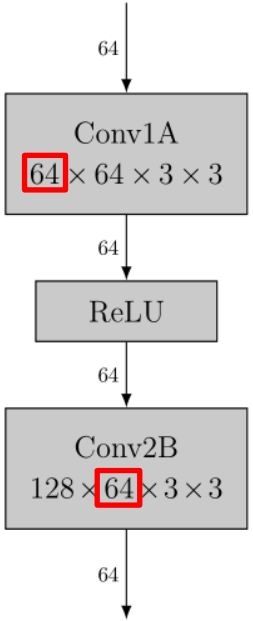

Challenges

α-schedule

- Type of decay (step, exponential, multiplicative, ...)
- Decay rate

Pretrained model

Should the weights of the pretrained model be freezed or updated?

Learning rate

- How big/small?
- Static or dynamic?
- Schedulers?

VGG16

Original performance:

- **Training accuracy: 100%**
- **Validation accuracy: 92%, Test accuracy: 91.3%**

VGG16

Original performance:

- **Training accuracy: 100%**
- **Validation accuracy: 92%, Test accuracy: 91.3%**

Standard pruning results:

- **Validation accuracy: 90%**
- **Pruned model size: ≈ 35% of original**

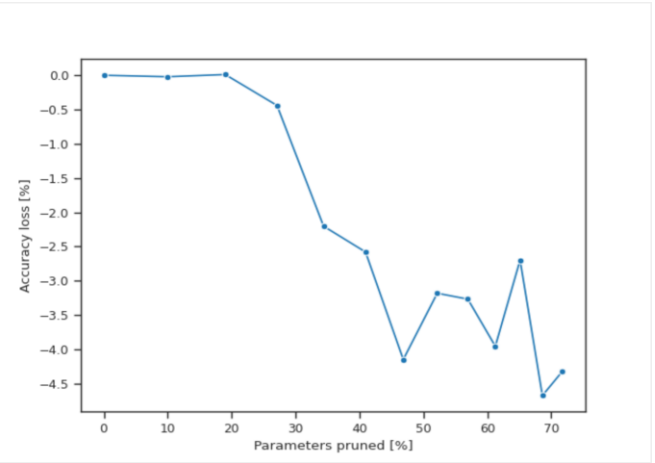

VGG16

Original performance:

- **Training accuracy: 100%**
- **Validation accuracy: 92%, Test accuracy: 91.3%**

Standard pruning results:

- **Validation accuracy: 90%**
- **Pruned model size: ≈ 35% of original**

⇒ Goal: **Prune 65% at once**

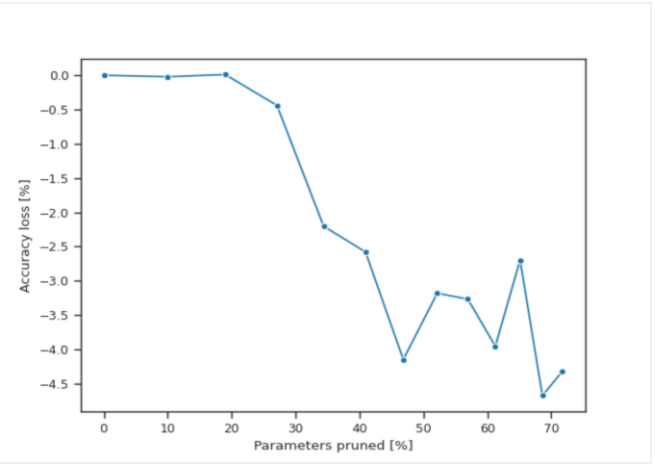

VGG16: Results

VGG16: Results

Failed attempts/mistakes:

- **Freeze pretrained + reduceLrOnPlateau-scheduler + constant initial LR (0.001)**
- **Step decay/Exp. decay**

VGG16: Results

Failed attempts/mistakes:

- **Freeze pretrained + reduceLrOnPlateau-scheduler + constant initial LR (0.001)**
- **Step decay/Exp. decay**

Final setup:

- **Update pretrained + step decay (-0.1) + reduceLrOnPlateau-scheduler**
- **At start: Initial LR of 0.001, at end: Repeated increase of initial LR x10**

VGG16: Results

Failed attempts/mistakes:

- **Freeze pretrained + reduceLrOnPlateau-scheduler + constant initial LR (0.001)**
- **Step decay/Exp. decay**

Final setup:

- **Update pretrained + step decay (-0.1) + reduceLrOnPlateau-scheduler**
- **At start: Initial LR of 0.001, at end: Repeated increase of initial LR x10**

Results:

- **Validation accuracy: 88.5%**
- **Parameters remaining: ≈ 35%**

	Original	Pruned
Time/img (GPU)	≈0.0531s	≈0.0529s
GPU VRAM	1.53GB	0.53GB
Time/img (CPU)	≈0.0409s	≈0.0257s
RAM consumption	3.64GB	2.32GB

RGB Autoencoder

Original performance (input of size 3x128x128, bottleneck of size 128x16x16):

- **Training loss: 1.14**
- **Validation loss: 1.40, Test loss: 1.41**

Original

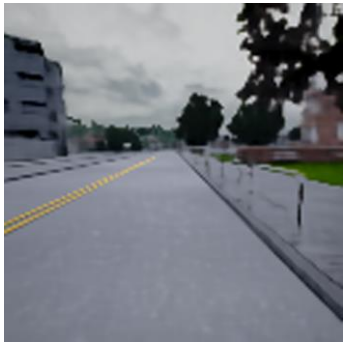

Reconstructed

Original

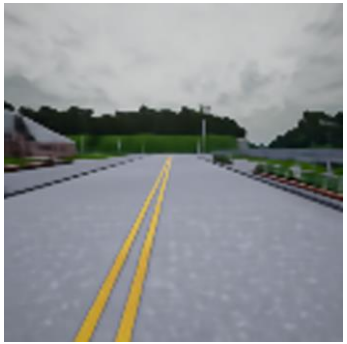

Reconstructed

RGB Autoencoder

Original performance (input of size 3x128x128, bottleneck of size 128x16x16):

- **Training loss: 1.14**
- **Validation loss: 1.40, Test loss: 1.41**

Original

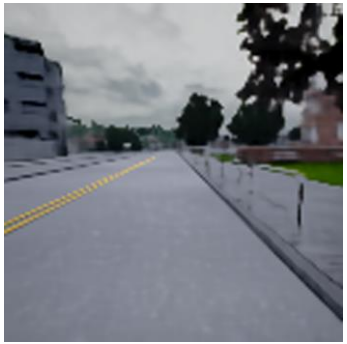

Reconstructed

Original

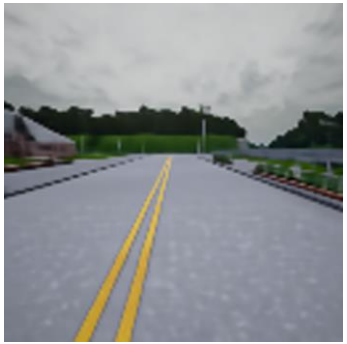

Reconstructed

⇨ Goal: **Prune 65% at once – bottleneck of size 44x16x16**

RGB Autoencoder: Approach and Results

RGB Autoencoder: Approach and Results

Possible approaches:

- **Prune encoder – then decoder**
- **Prune en- and decoder simultaneously**

RGB Autoencoder: Approach and Results

Possible approaches:

- **Prune encoder – then decoder**
- **Prune en- and decoder simultaneously**

Setup: **Same setup as with VGG**

RGB Autoencoder: Approach and Results

Possible approaches:

- **Prune encoder – then decoder**
- **Prune en- and decoder simultaneously**

Setup: **Same setup as with VGG**

Results – Prune encoder:

- **Training loss: 2.31**
- **Validation loss: 2.59**
- **Test loss: 2.65**

	Original	Pruned
Time/img (GPU)	≈0.0573s	≈0.0582s
GPU VRAM	6.83GB	5.23GB
Time/img (CPU)	≈0.0891s	≈0.0836s
RAM consumption	7.97GB	7.89GB

RGB Autoencoder: Results – Prune Encoder

Img. 1

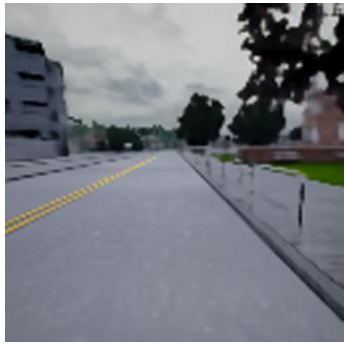

Img. 2

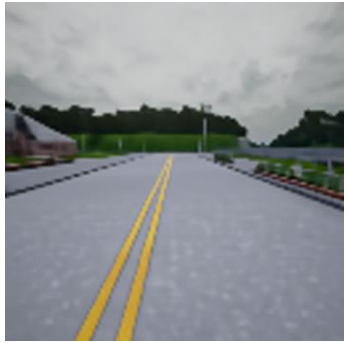

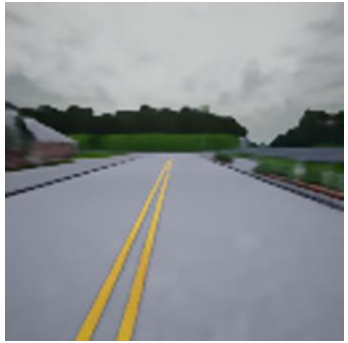

Original

Reconstructed
(before pruning)

Reconstructed
(after pruning,
before multiplication)

Reconstructed
(after pruning,
after multiplication)

RGB Autoencoder: Intermediate Results – Prune Both

Img. 1

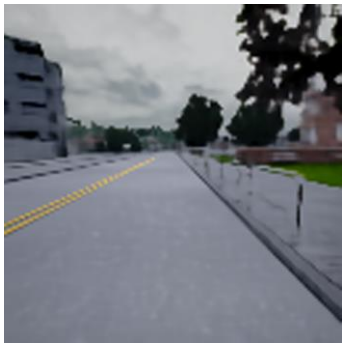

Img. 2

Original

Reconstructed
(before pruning)

Reconstructed
(after pruning,
before multiplication)

Conclusion

α-schedule

- Type of decay (step, exponential, multiplicative, ...)
- Decay rate

Pretrained model

Should the weights of the pretrained model be freezed or updated?

Learning rate

- How big/small?
- Static or dynamic?
- Schedulers?

Conclusion

α-schedule

- Type of decay (step, exponential, multiplicative, ...)
- Decay rate

Pretrained model

Should the weights of the pretrained model be freezed or updated?

Learning rate

- How big/small?
- Static or dynamic?
- Schedulers?

Open question

Conclusion

α-schedule

- Type of decay (step, exponential, multiplicative, ...)
- Decay rate

Open question

Pretrained model

Should the weights of the pretrained model be freezed or updated?

Update weights

Learning rate

- How big/small?
- Static or dynamic?
- Schedulers?

Conclusion

α-schedule

- Type of decay (step, exponential, multiplicative, ...)
- Decay rate

Open question

Pretrained model

Should the weights of the pretrained model be freezed or updated?

Update weights

Learning rate

- How big/small?
- Static or dynamic?
- Schedulers?

Initial LR of 0.001
Increase at local minima

Future Work

- Prune autoencoder decoder
- Compare results of both autoencoder approaches

- …

- Prune autoencoder with skip connections
- In general: Test new pruning approach with a larger α-decay

- …

- Goal: Prune DCGAN

Thanks!

Does anyone have any questions?

Sources

Presentation template by Slidesgo

Original model

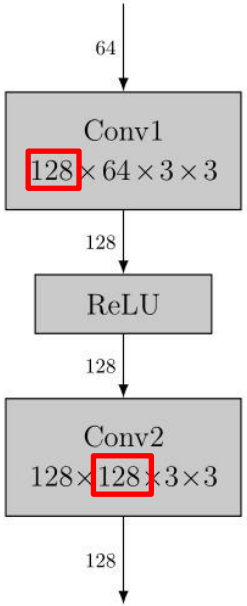

With parallel branches

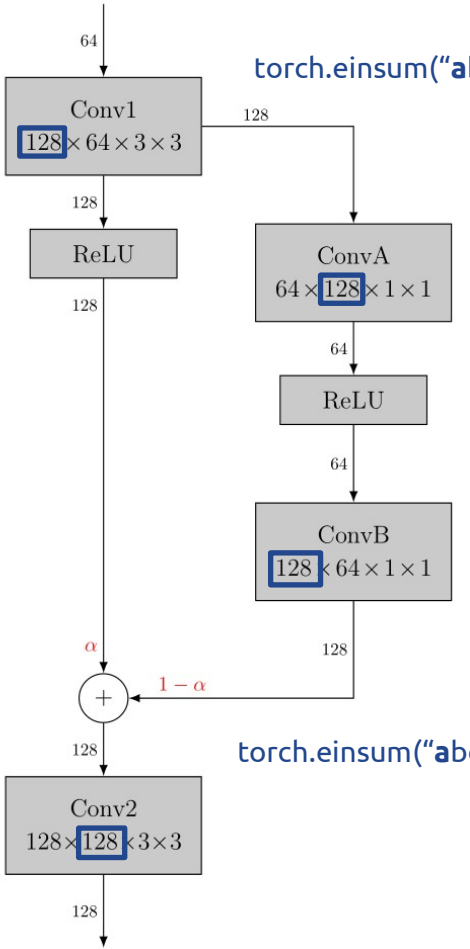

torch.einsum("**a**bcd, e**a**fg -> ebcd", conv1_weights, convA_weights)

If bias:
W2(W1x + b1) + b2 = (W2W1)x + (W2b1 + b2)

New weights New bias

torch.einsum("**a**bcd, e**a**fg -> eb**f**g", convB_weights, conv2_weights)

Original model

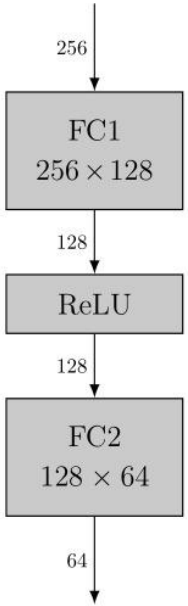

With parallel branches

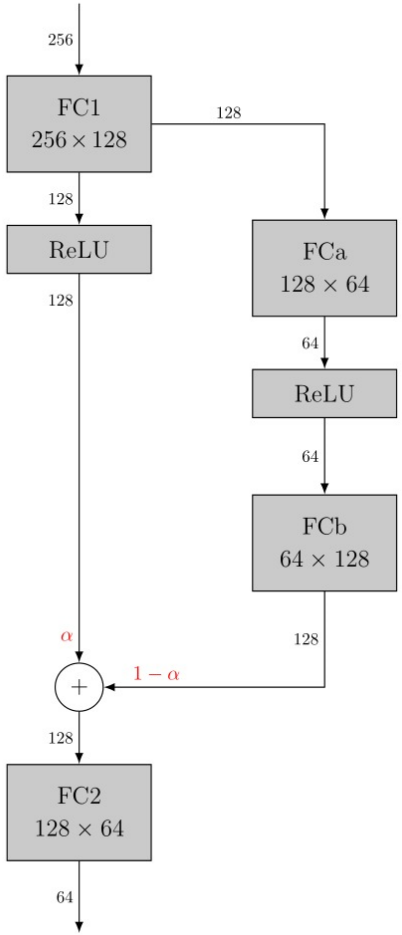

Pruned model

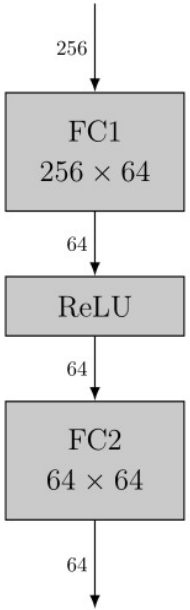

Original model

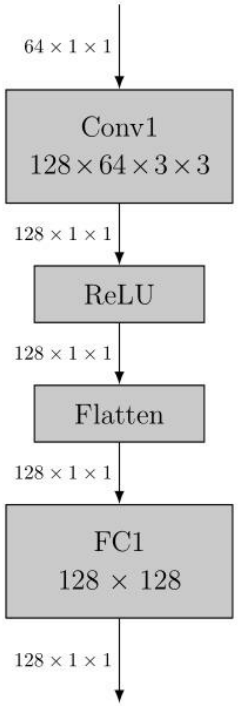

With parallel branches

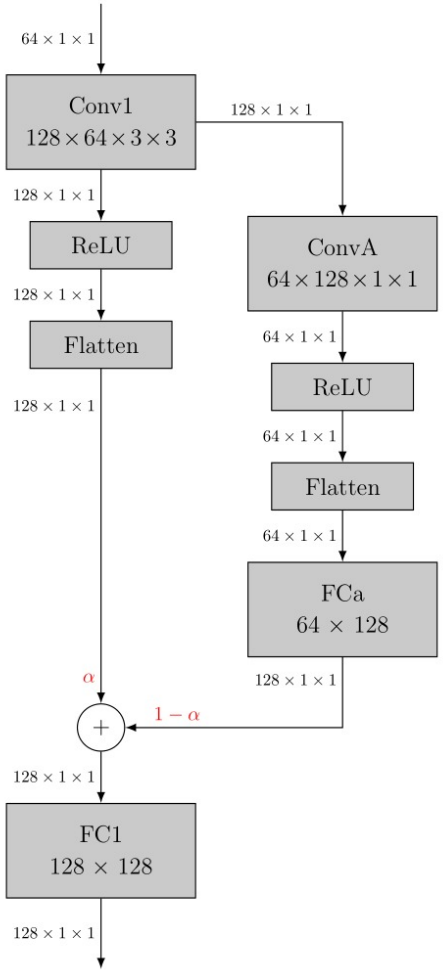

Pruned model

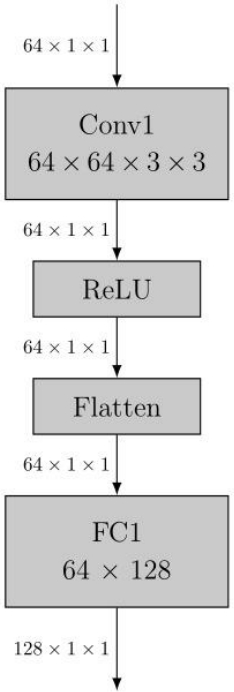

Original model

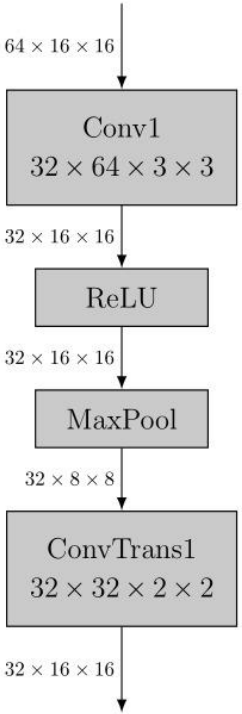

With parallel branches

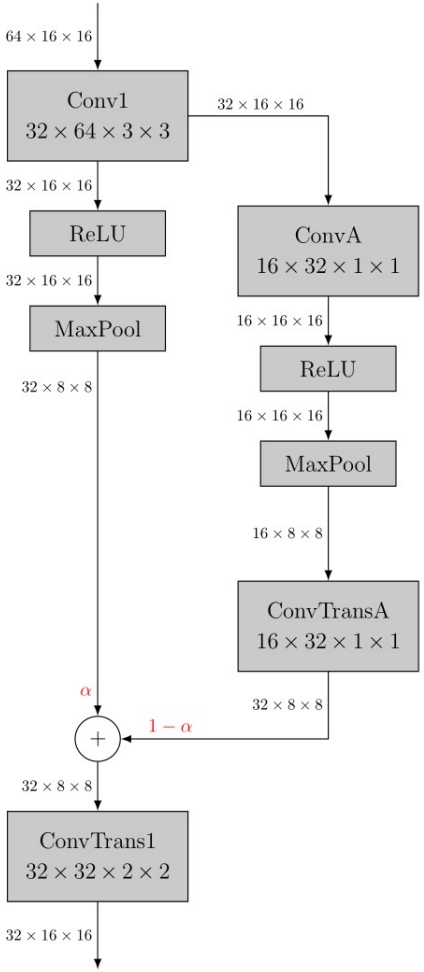

Pruned model